

RESEARCH METHODOLOGY

Prompts That Genuinely Changed the Architecture

Eight exchanges, pulled from the actual build conversation across Phases 1–3, each one tied to a real decision it changed and why that decision mattered.

OC Venue Activation Knowledge Base · August 2026

Why This List Exists

The task brief's guidance on using AI as a tutor rather than an oracle asks for a record of "the prompts that genuinely changed your decisions" — not a transcript, and not every message, just the ones that actually reversed or forced a real technical choice. What follows are the eight that qualify, in the order they happened.

"Why do you recommend the OpenAI text-embedding?"

Pushed for a justification instead of accepting the default. The follow-up questions this opened — cost, and whether Supabase's built-in option was genuinely free and reversible — led to choosing **gte-small running locally in a Supabase Edge Function** over a paid third-party API: zero cost, zero account setup, and a confirmed low-risk choice since switching later only means re-embedding existing chunks.

"Will this be an issue if I scale to Orange County?"

Produced the single most important architectural principle in the project: **city disambiguation has to be a metadata filter, not an embedding-model choice**. Nine cities' ordinances are structurally near-identical, so similarity search alone risks matching the right kind of fact to the wrong city. The fix — filter by city before ranking by similarity — is written directly into the Phase 4 plan and saved to persistent memory so it isn't lost before Phase 4 starts.

"How do you know it's a weak match? It shows a 77.8% similarity."

Refused to accept an assertion at face value. This forced an actual empirical test — a nonsense control query run through the same search — which revealed the corpus's real similarity floor (~0.74, versus ~0.90 for genuinely relevant matches). That number later explained *why* two different out-of-scope test questions failed for two different reasons.

"Is Yelp Fusion free? ... I'm seeing different plans." (with a screenshot)

Overtaken a stale assumption that Yelp was a simple free starting point. It isn't anymore. This caused the pivot to Google Places, and more importantly established a pattern for the rest of the build: verify current pricing and terms before committing, rather than trust training data. That habit caught the same staleness again later, more than once.

"Is there anything about the instructions that calls for Railway?" → "I'm referring to the 'Build a Knowledge Base' document."

The biggest correction in the project. This sent the process back to actually read the original task brief PDF for the first time — which names Railway as a required deliverable three separate times. Without this question, deployment would likely have moved to Vercel, satisfying the spirit of "a live app" but not the letter of the actual assignment.

"Can we begin working on Phase 2 without destroying the Phase 1 demo?"

Surfaced a real gap: despite substantial Phase 1 work, nothing had actually been committed to version control yet. Produced the git-tagging discipline used for the rest of the build — a permanent snapshot at the end of every phase, and real commits before starting new work.

“What does it mean that I can ‘check out and redeploy an old snapshot on demand?’ ... I need to show my work.”

Pushed past a vague technical explanation to name the actual concern: durable, presentable proof, not just technical recoverability. This settled the deployment strategy — one evolving live URL plus permanent tags in version control, rather than the cost and upkeep of running a separate live environment per phase.

“I would like to include pricing information for the venues... What was the actual purpose of including all those venues?”

Direct, well-founded pushback on a Phase 2 deliverable. Led to actually building the market-rate benchmarks table that had been in the data model since Phase 1 planning but never implemented, and sourcing real venue pricing to close a gap the evaluation had already flagged but nothing had acted on yet.

The Pattern

Almost none of these were "explain this to me" questions. They were either pushback on an unverified claim — *how do you know, is that actually free, does anything actually require this* — or a why/what-for question that exposed a decision made without a clearly articulated reason. Both were consistently more useful for shaping the architecture than open-ended requests to keep going.